
A Learning-Augmented Approach to Online Allocation Problems

Ilan Reuven Cohen

Faculty of Engineering
Bar-Ilan University
Ramat Gan, Israel.

ilan-reuven.cohen@biu.ac.il

Debmalya Panigrahi

Department of Computer Science
Duke University
Durham, NC, USA.
debmalya@cs.duke.edu

Abstract

In online allocation problems, an algorithm must choose from a set of options at each step, where each option incurs a set of costs/rewards associated with a set of d agents. The goal is to minimize/maximize a function of the accumulated costs/rewards assigned to the agents over the course of the entire allocation process. Such problems are common in combinatorial optimization, including minimization problems such as machine scheduling and network routing, as well as maximization problems such as fair allocation for welfare maximization.

In this paper, we develop a general learning-augmented algorithmic framework for online allocation problems that produces a nearly optimal solution using only a single d -dimensional vector of learned weights. Using this general framework, we derive learning-augmented online algorithms for a broad range of application problems in routing, scheduling, and fair allocation. Our main tool is convex programming duality, which may also have further implications for learning-augmented algorithms in the future.

1 Introduction

In the last few years, tremendous advances in machine learning have triggered much interest and rapid progress in a new domain called *learning-augmented* online algorithms. In many practical scenarios, one needs to solve optimization problems with dynamically evolving requirements (called online algorithms), where uncertainty about the future significantly affects the quality of the solution. The basic premise of the learning-augmented paradigm is to use (possibly noisy) machine-learned predictions as a proxy for the future, and design online algorithms that can take advantage of good predictions while not falling prey to bad ones. (The resulting algorithms are sometimes also referred to as “online algorithms with predictions”.) The explosive growth of this new area over the last few years has spanned a large variety of domains such as caching, scheduling, matching, clustering, network design, and many others (see related work), and generated much excitement at the intersection of algorithm design and machine learning. While many novel and clever algorithms have been designed in the learning-augmented setting, these are typically tailored to the requirements of a specific problem at hand, and do not generalize to broader problem classes. Indeed, there is a surprising absence of unifying algorithmic models and general-purpose tools for online algorithms with predictions, particularly when contrasted with other models of computation where algorithmic progress has largely relied on broad, unifying models and algorithmic techniques.

In this paper, we aim to partially rectify this situation by studying a broad class of problems called online allocation in the learning-augmented setting. The setup is the following: in each online step, the algorithm is presented a set of options and has to choose one of them. Each option is represented by a vector, and the final objective is a (linear or non-linear) function of the (coordinate-wise) sum

of all the chosen vectors. As we will soon see, this captures a broad range of problems in network routing, scheduling, fair allocation, etc. In this paper, we develop a general algorithmic strategy for this entire class of problems in the learning-augmented setting that obtains a nearly optimal (fractional) solution. This is in sharp contrast to strong (super-constant) lower bounds without ML advice, for almost all problems captured by this framework.

The online allocation problem that we described above has many applications. Among the most important is that of online routing. In this problem, we are given a network and in each online step, the algorithm has to choose a route from a given source to a destination node. Eventually, the goal is to minimize the maximum congestion of any network link, which is defined as the number of chosen routes using that link. This is a classical and well-studied problem in online algorithms, and it is well-known that the best competitive ratio¹ achievable is logarithmic in the size of the network [3] (see also [4, 6, 7]). In sharp contrast, we obtain a $(1 + \epsilon)$ -approximation for any $\epsilon > 0$ using learning augmentation, by modeling the problem as online allocation where each coordinate in the vector represents a network link and the goal is to minimize the ℓ_∞ -norm of the congestion vector.

A second application domain for online allocation is that of online scheduling problems. Consider the problem of processing jobs online where each job comes with a slate of options given by vectors that represent the processing time on different machines. Overall, the goal is to minimize the makespan, i.e., the maximum load on any machine. This problem is also described via online allocation, where the coordinates of the vectors represent machines and the objective is again the ℓ_∞ -norm of machine loads given by the sum of the selected vectors. Our online allocation algorithm gives a $(1 + \epsilon)$ -competitive learning-augmented solution to this problem, which again sharply contrasts with lower bounds that are logarithmic in the number of machines in the absence of learning augmentation [10].

We note that the flexibility of the online allocation framework also allows us to extend these problem formulations immediately to a much broader set of objectives, such as all ℓ_p -norm objectives [5] rather than just the ℓ_∞ -norm. Another interesting direction is that of maximization problems. In contrast to minimization problems where the vector coordinates can be thought of as resources whose use needs to be minimized, we now imagine the coordinates to represent value added to agents when a particular option is selected. This allows us to model fair allocation problems, where a set of items have to be assigned in a way that maximizes the minimum value along all agents (or another function such as Nash social welfare, p -means, etc.) [13, 15] Again, we observe the sharp contrast between the traditional online setting without learning augmentation, where strong polynomial lower bounds exist for these problems [34, 18] and the learning-augmented setting where our online allocation algorithm immediately yields a $(1 - \epsilon)$ -approximation for any $\epsilon > 0$.

We end this section by commenting on the machine-learned parameters that we use in our learning-augmented algorithm for the online allocation problem. As described above, we show that there exists a set of learned parameters that allows us to approximate the optimal objective within a $(1 \pm \epsilon)$ factor of the optimal value. Importantly, we show that these parameters are bounded as a function of ϵ and the number of coordinates (rather than the number of steps or options) and, moreover, that our learning-augmented algorithm is resilient to small errors in the values of the learned parameters. This allows us to derive PAC-learning results for the learned parameters and to use them in the learning-augmented algorithm to obtain a $(1 \pm O(\epsilon))$ -approximation.

Organization. The formal model and statement of results appear in Section 2. Section 3 presents related work. The main result, which provides the learning-augmented algorithm, appears in Section 4. Applications of this result to routing and fair allocation problems are discussed in Section 5. In Appendix C, we derive sample complexity bounds for the learned parameters in the PAC model. All omitted details of proofs are in the appendix.

2 Our Contributions

First, we define the online allocation framework:

Online Allocation Problem. In this problem, we receive an initial input of a function f over a bounded domain in \mathbb{R}_+^d . Subsequently, we proceed in steps. At every time step $t = 1, \dots, T$, we receive a set $A_t \subseteq \mathbb{R}_+^d$ of d -dimensional vectors. We must select one vector $\mathbf{v}_t^\dagger \in A_t$ before

¹As is usual in online algorithms, performance of an algorithm is measured via its competitive ratio, which is the worst-case ratio of the algorithm's objective to that of an optimal solution.

proceeding to time step $t + 1$, using only information available up to time t . Denote the result vector $\mathbf{v}_{\text{tot}}^\dagger := \sum_{t=1}^T \mathbf{v}_t^\dagger$. The goal is to maximize or minimize $f(\mathbf{v}_{\text{tot}}^\dagger)$ for some objective function $f(\cdot)$.

For ease of notation, we fix an arbitrary ordering of the vectors in A_t , and denote by $v_{t,k} \in \mathbb{R}_+^d$ the k th vector in A_t , i.e., $A_t = \langle v_{t,1}, \dots, v_{t,|A_t|} \rangle$. Accordingly, we define $K(t) = \{1, \dots, |A_t|\}$ as the set of indices of these vectors, and for each $k \in K(t)$, we let $v_{i,t,k}$ denote the i th coordinate of vector $v_{t,k}$ in A_t . Our results apply to *well-behaved functions* $f(\cdot)$ that we define as follows:

Well-Behaved Functions. Let $f : \mathbb{R}_+^d \rightarrow \mathbb{R}_+$ be the objective function defined on the result vector. Then, f is well-behaved if it satisfies the following properties:

Monotonicity: f is said to be *monotone* if for any $\ell, \ell' \in \mathbb{R}_+^d$ such that $\ell_i \geq \ell'_i$ for all $i \in [d]$, we have $f(\ell) \geq f(\ell')$.

Homogeneity: f is said to be *homogeneous* if for any $\ell, \ell' \in \mathbb{R}_+^d$ such that $\ell'_i = \alpha \cdot \ell_i$ for all $i \in [d]$, we have $f(\ell') = \alpha \cdot f(\ell)$.

These properties are satisfied by most objective functions studied in the literature including linear functions, p -norms, Nash Social Welfare (which is the geometric mean), p -means, etc.

Our algorithmic scheme will focus on developing a fractional selection rule, which can be interpreted as a probability distribution over the different choices in every step. We formally describe a fractional solution as follows:

Fractional Solution. At step t , the algorithm has to assign a fractional value to each option $x_{t,k} \in [0, 1]$ for $k \in K(t)$ such that $\sum_{k=1}^{|K(t)|} x_{t,k} = 1$. Accordingly, we have

$$\mathbf{v}_{\text{tot}}^f := \sum_{t=1}^T \sum_{k=1}^{|A_t|} \mathbf{v}_{t,k} \cdot x_{t,k}.$$

Note that in some scenarios such as the allocation of divisible items or network flow routing, a fractional solution is already sufficient.

Learning-Based Online Scheme for Allocation Problems. Our main result is that for any d -dimensional instance of the online allocation problem with a well-behaved objective function, there exists a vector $\alpha \in \mathbb{R}^d$ such that an online algorithm guided by α achieves a competitive ratio of $1 + \epsilon$ for minimization or $1 - \epsilon$ for maximization, for any $\epsilon > 0$.

Theorem 2.1. *Given an instance of the online allocation problem with a well-behaved objective and any $\epsilon > 0$, there exists a set of learned parameters $\alpha \in \mathbb{R}^d$ and an online algorithm that uses α such that the resulting fractional solution achieves a $(1 \pm O(\epsilon))$ -approximation.*

The online algorithm that establishes this result uses a simple *exponential assignment rule*, where given a learned parameter vector $\alpha \in \mathbb{R}^d$, the fractional allocation $x_{tk}(\alpha)$ of option $k \in K(t)$ is:

$$x_{tk}(\alpha) \propto \exp \left(- \sum_{i=1}^d \alpha_i \cdot \frac{v_{i,t,k}}{\mathbf{m}_t(v)} \right), \text{ where } \mathbf{m}_t(v) = \min \{v_{itk} | v_{itk} > 0, i \in [d], k \in K(t)\}.$$

We further show that the learned parameters can be bounded as a function of d and ϵ , and that the algorithm is robust to small perturbations in the parameter values. This structure allows the parameter space to be both bounded and discretized in terms of d, ϵ , independent of the time horizon. This property is crucial in applications where the number of steps (which may represent flow requests, jobs, or items) is significantly larger than the number of dimensions (which may represent network links, machines, or agents). Moreover, the discretization of the parameter space ensures that the parameters are learnable, and allows us to bound the number of samples required for learning.

Learnability of Parameters. Following the PAC framework of [26], we establish the learnability of parameter vectors for well-behaved objectives, under the additional assumption of *superadditivity* for maximization or *subadditivity* for minimization, as introduced in [18].

We consider a setting in which each online step (i.e., a set of vectors) is drawn independently (though not necessarily identically) from a distribution, and each step contributes only a small fraction of the total allocation. This latter property is formalized by the standard *small items assumption* (see, e.g., [1, 26, 24, 29]), which we state below:

Small Items Assumption: There exists a $\zeta = \Theta\left(\frac{\log d}{\epsilon^2}\right)$ such that $v_{itk} \leq \frac{L}{\zeta}$ for every $i \in [d]$, $t \in [T]$, and $k \in K(t)$.

Under this assumption, we derive a sample complexity bound for approximately learning a parameter vector α that ensures near-optimal performance. The following informal theorem summarizes our main learnability result; see Appendix C for details:

Theorem 2.2. *Under standard PAC assumptions, the small items assumption, and mild regularity conditions on the objective (subadditivity for minimization and superadditivity for maximization), the parameter vector $\alpha \in \mathbb{R}^d$ can be learned from $O\left(\frac{d}{\log d} \cdot \log \frac{d}{\epsilon}\right)$ i.i.d. samples.*

Rounding the Fractional Solution. As mentioned earlier, a fractional solution is already sufficient for many applications. Furthermore, we show that under the small items assumption (see above), a fractional solution can be converted into a randomized online integral solution with a $(1 \pm \epsilon)$ loss in the objective value, where the parameter ϵ only depends on the parameters of the small items assumption.

Consider the following online randomized rounding algorithm: given an online fractional solution x_{tk} for $t \in [T]$, the algorithm chooses option k with probability $x_{t,k}$. Under the small items assumption, standard Chernoff bounds ensure that with high probability, the maximum coordinate in the rounded vector remains within a multiplicative factor of $(1 \pm O(\epsilon))$ of its fractional counterpart. We get the following theorem:

Theorem 2.3. *Let \mathbf{v}_{tot}^f be a fractional solution, and let \mathbf{v}_{tot}^\dagger be the integer solution produced by randomized rounding of this fractional solution. Then, under the small items assumption, the following holds with high probability:*

$$\begin{aligned} \max_{i \in [d]} \mathbf{v}_{tot,i}^\dagger &\leq (1 + \epsilon) \cdot \max_{i \in [d]} \mathbf{v}_{tot,i}^f, \text{ and} \\ \min_{i \in [d]} \mathbf{v}_{tot,i}^\dagger &\geq (1 - \epsilon) \cdot \min_{i \in [d]} \mathbf{v}_{tot,i}^f. \end{aligned}$$

Now, combining the above theorem with Theorem 2.1, we get the following result for integer solutions under the small items assumption:

Corollary 2.4. *Given an instance of the online allocation problem with a well-behaved objective and under the small-item assumption, for any $\epsilon > 0$, there exists a set of learned parameters $\alpha \in \mathbb{R}^d$ and an online algorithm that uses α such that the resulting integral solution achieves a $(1 \pm O(\epsilon))$ -approximation.*

Robustness–Consistency Tradeoff. Ideally, a learning-augmented algorithm should simultaneously exploit accurate predictions to achieve near-optimal performance (consistency), while also maintaining strong worst-case guarantees when predictions are unreliable (robustness). The (not learning-augmented) worst case lower bounds for allocation problems with minimization (resp., maximization) objectives is $\Omega(\log d)$ (resp., $\Omega(d)$). Formally, an algorithm is said to be γ -consistent and δ -robust if it achieves a γ -approximation under accurate predictions (consistency), and a δ -approximation in the worst case when predictions are unreliable (robustness). In Appendix D, we give a slight modification of our scheme that preserves robustness while achieving consistency for both minimization and maximization objectives.

Theorem 2.5. *For minimization objectives, there exists an algorithm that uses $\alpha \in \mathbb{R}^d$, a predicted parameter vector, and achieves 1-consistent and $O(\log d)$ -robust approximation.*

For maximization objectives, there exists an algorithm that uses $\alpha \in \mathbb{R}^d$, a predicted parameter vector, and a parameter λ , which achieves a $(1 - \lambda)$ -consistent and $\lambda \cdot d$ -robust approximation.

Applications. In Section 5, we illustrate the utility of our framework in handling a broad class of objectives by demonstrating its use in two applications. The first is for the online routing problem, where we are given a network and in each time-step, there is request to route a given value of flow from a source to a destination vertex [3]. The flow needs to be routed in a way that minimizes maximum congestion on any edge cumulatively across all time-steps. We give the first learning-augmented algorithm for this problem as a simple corollary of our general online allocation framework. Next,

we consider a maximization problem, that of allocating items to agents so as to maximize the Nash Social Welfare [13]. Again, we give a learning-augmented algorithm for this problem as a simple corollary of our general framework, matching previous results in [18]. We note that our framework applies to many other application problems in domains such as scheduling (makespan and ℓ_p -norm minimization) and fair allocation (Santa Claus and p -means maximization) that we do not state here for brevity.

3 Related Work

Previous work that is mostly closely related to ours is on learning-augmented online scheduling and assignment problems [23, 24, 26, 18]. These works focus on the assignment of items that arrive online to agents, where each choice affects only a single agent. In our setting, this corresponds to each choice being represented by a d -dimensional vector with a single nonzero dimension, where d corresponds to the number of agents.

Lattanzi et al. [23] focused on minimizing the makespan (maximum load) for restricted assignment. They showed that a suitable set of learned parameters enables a proportional allocation rule that obtains a nearly optimal result. In [24], it was further shown that these parameters are PAC-learnable.

Li and Xian [26] generalized this framework to handle unrelated scheduling. Their scheme introduces two learned parameters per machine. The first set of parameters transform the instance into a restricted-related setting, and they show that proportional assignment using the second set of parameters then applies to this restricted setting.

Cohen and Panigrahi [18] improved this result by using only a single set of parameters in an exponential allocation scheme, and also extended it to handle a large number of maximization and minimization objectives.

Our scheme generalizes these previous works to accommodate general vectors, which is crucial for applications such as online routing that cannot be solved using the previous techniques. Furthermore, we simplify the previous approaches as well by requiring only a single variable per dimension, thereby eliminating the dependence on an additional exponent base which further simplifies the analysis of the existence of such parameters. Our work introduces a new set of techniques based on perturbation/sensitivity analysis because the ideas previously used for bounding the learned parameters do not generalize to the case of arbitrary vector options.

In [14], the authors study the Online Nash Social Welfare problem with predictions. The goal is to compute an online divisible allocation of goods among d agents in a way that balances fairness and efficiency. In their setting, the prediction for each agent is their *monopolist value*—the utility the agent would obtain if all resources were allocated solely to them—and their algorithm achieves an $O(\log d)$ -approximation under this assumption. In contrast, we show that by learning a different d -dimensional parameter vector, it is possible to achieve a nearly optimal allocation.

Various papers [12, 35, 19, 20] investigate the robustness–consistency tradeoff in the context of general packing and covering problems. These algorithms typically introduce an additional parameter $\lambda \in [0, 1]$ to encode the algorithm’s confidence in the prediction. A primal-dual scheme is then employed to interpolate between a worst-case baseline and a prediction-driven strategy. Unlike our approach, these methods assume that the prediction consists of the full solution, which is often impractical in real-world settings.

More broadly, the study of learning-augmented online algorithms was initiated by Lykouris and Vassilvitskii [27] in the context of the caching problem and has since grown into a prominent research area. This framework enhances online algorithms by incorporating machine-learned predictions about the future, allowing them to surpass pessimistic worst-case competitive bounds. Over the past few years, numerous online allocation problems have been explored within this framework, including applications in scheduling [33, 8, 9, 11, 21, 30], online matching [2, 17, 22], and ad delivery [28, 25]. For a broader overview of learning-augmented online algorithms, we refer the reader to the surveys by Mitzenmacher and Vassilvitskii [31, 32].

4 Online Allocation for a Well-Behaved Objective via Learned Parameters

In this section, we show that for the general case of the online allocation problem, there exists a set of learned parameters that guarantees a near-optimal solution. Specifically, we show the following theorem, which is a more refined version of (and establishes) Theorem 2.1:

Theorem 4.1. *Given an instance of the online allocation problem with a well-behaved objective and any $\epsilon > 0$, there exists a set of learned parameters $\alpha \in \text{NET}(q, s)$ and an online algorithm that uses α such that the resulting fractional solution achieves a $(1 \pm O(\epsilon))$ -approximation.*

Here, $\text{NET}(q, s) = \left\{ \frac{i}{S} \mid i \in [0, q \cdot s] \right\}^d$ denotes a d -dimensional discrete net with parameters q, s that are bounded in $\text{poly}(d, 1/\epsilon)$.

For the minimization of a well-behaved function, we first consider the **MinMax** objective, which is defined as the minimization of f , where $f(v) = \max_{i \in [d]} v_i$. Similarly, for the maximization of a well-behaved function, we first consider the **MaxMin** objective, which is defined as the maximization of f , where $f(v) = \min_{i \in [d]} v_i$.

In this section, we focus on the **MinMax** objective. In Appendix E, we address the **MaxMin** objective. In Appendix F, we explain how to extend these results to general *well-behaved* objectives.

4.1 Learning-Augmented Online Allocation for the MinMax Objective

We now state the main result for the **MinMax** objective:

Theorem 4.2. *Given an instance of the online allocation problem with a **MinMax** objective and any $\epsilon > 0$, there exists a set of learned parameters $\alpha \in \text{NET}\left(\frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right), \frac{d^3}{\epsilon^3}\right)$ and an online algorithm that uses α such that the resulting fractional solution is a $(1 + O(\epsilon))$ -approximation.*

To prove this result, we first apply a pre-processing step to convert arbitrary instances of the problem to structured instances that we say are *balanced*. Note that the conversion to balanced instances is an algorithmic technique and not a restriction on the input. We then leverage convex programming duality on a max-entropy style convex programming formulation to show the existence of learned parameters that can be used to obtain a near-optimal solution online. At this stage, the learned parameters can be of arbitrary precision and are not necessarily efficiently learnable. To ensure the latter, we need to obtain parameters that belong to $\text{NET}\left(\frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right), \frac{d^3}{\epsilon^3}\right)$. We do this in the last step by using tools from perturbation and sensitivity analysis in the convex programming literature. Together, these steps establish Theorem 4.2.

Preprocessing of MinMax Instances. Recall that the exponential assignment scheme computes online allocations based on the ratio $\frac{v_{i,t,k}}{\mathbf{m}_t(v)}$. This ratio can be highly sensitive and potentially unbounded under small perturbations to the input, thereby making it impossible to execute the last step of designing learned parameters of bounded precision. To overcome this difficulty, we introduce a pre-processing step that transforms the original instance into a *balanced* instance, incurring only a small loss in the objective value.

An instance $\tilde{I}(\tilde{v}, \tilde{K})$ is said to be *balanced* if it satisfies the following condition for all $t \in [T]$, $i \in [d]$, and $k \in K(t)$:

$$\frac{\tilde{v}_{itk}}{\mathbf{m}_t(\tilde{v})} \in \{0\} \cup \left\{ (1 + \epsilon)^b \mid b \in \mathbb{Z}_+ \cap \left[0, \log_{1+\epsilon} \left(\frac{d^2}{\epsilon^2} \right) \right] \right\}.$$

Note that in a balanced instance, we have $\frac{v_{i,t,k}}{\mathbf{m}_t(v)} \leq \frac{d^2}{\epsilon^2}$. Moreover, by omitting redundant options, we may assume that the number of options per step is bounded, i.e., $\ln |K(t)| = O(d \cdot \ln(d/\epsilon))$. We construct a balanced instance $I(\tilde{v}, \tilde{K})$ as follows:

Given an instance $I(v, K)$,

1. Define \hat{v} such that

$$\hat{v}_{itk} = \begin{cases} 0 & \text{if } \frac{v_{itk}}{\max_{i' \neq i} v_{i'tk}} < \frac{\epsilon}{d} \\ v_{itk} & \text{otherwise.} \end{cases}$$

2. Define \tilde{K} by removing option k' from step t whenever there exists another option k with

$$\max_i \hat{v}_{i,t,k'} > \frac{d}{\epsilon} \cdot \max_i \hat{v}_{i,t,k}.$$

3. Define \tilde{v} as follows: if $\hat{v}_{itk} = 0$, then set $\tilde{v}_{itk} = 0$; otherwise, set

$$\tilde{v}_{itk} = \mathbf{m}_t(\hat{v}) \cdot (1 + \epsilon)^{\lfloor \log_{1+\epsilon} \left(\frac{\hat{v}_{itk}}{\mathbf{m}_t(\hat{v})} \right) \rfloor}.$$

By definition, $I(\tilde{v}, \tilde{K})$ is a balanced instance. In Appendix A, we show that transforming $I(v, K)$ into $I(\tilde{v}, \tilde{K})$ incurs only a $(1 + O(\epsilon))$ loss in the objective value.

Lemma 4.3. *Let $I(v, K)$ be an instance of the allocation problem with the **MinMax** objective, and let $\epsilon > 0$. Then, any $(1 + \epsilon)$ -approximate solution to the balanced instance $I(\tilde{v}, \tilde{K})$ constructed via the pre-processing algorithm yields a $(1 + O(\epsilon))$ -approximate solution to $I(v, K)$.*

Learned Parameters for Balanced Instances. In light of Lemma 4.3, it suffices to only consider balanced instances in Theorem 4.2. We restate this goal as follows:

Lemma 4.4. *Given a balanced instance of the online allocation problem with a **MinMax** objective and $\epsilon > 0$, there exists a set of parameters $\alpha \in \text{NET} \left(\frac{d^2}{\epsilon} \cdot \ln \left(\frac{d}{\epsilon} \right), \frac{d^3}{\epsilon^3} \right)$ such that the fractional solution defined by the exponential assignment scheme with parameters α is a $(1 + O(\epsilon))$ -approximation to the optimal objective.*

The remainder of this section gives a proof of Lemma 4.4. We do this in two steps. First, we prove the existence of parameters $\alpha^{(\epsilon)}$ that are bounded as a function of d, ϵ using which the exponential assignment scheme produces an $(1 + \epsilon)$ -approximate optimal solution. (For the special case of $\epsilon = 0$, the solution is precisely optimal but the parameters may be unbounded.) Then, we show that small errors (up to ϵ^3/d^3) in the parameter values incur only a $(1 + O(\epsilon))$ loss in the objective value, enabling discretization of the learned parameters to $\text{NET} \left(\frac{d^2}{\epsilon} \cdot \ln \left(\frac{d}{\epsilon} \right), \frac{d^3}{\epsilon^3} \right)$.

We use L^* to denote the optimal **MinMax** value, and define the following convex program for $\epsilon > 0$:

$$\begin{aligned} \min \quad & \sum_{t \in [T]} \mathbf{m}_t(v) \sum_{k \in K(t)} x_{tk} \cdot (\ln x_{tk} - 1) \\ \text{s.t.} \quad & \sum_{t \in [T]} \sum_{k \in K(t)} v_{itk} \cdot x_{tk} \leq L^* \cdot (1 + \epsilon), \quad \forall i \in [d], \\ & \sum_{k \in K(t)} x_{tk} = 1, \quad \forall t \in [T], \\ & x_{tk} \geq 0, \quad \forall k \in K(t), t \in [T] \end{aligned}$$

Figure 1: Convex Programming Formulation for the **MinMax** Objective

Lemma 4.5. *Given an instance of the online allocation problem with the **MinMax** objective and any $\epsilon \geq 0$, there exists a vector $\alpha^{(\epsilon)} \in \mathbb{R}_+^d$ such that the fractional solution defined by the exponential assignment scheme with parameters $\alpha^{(\epsilon)}$ is $(1 + O(\epsilon))$ -approximately optimal.*

Proof. Given such instance and for fixed ϵ consider the convex program of Figure 1. By our assumption, L^* is the optimal **MinMax** objective therefore there exists a feasible solution for the convex program for any $\epsilon \geq 0$. Accordingly, define the Lagrangian $L(x, \alpha, \beta)$ as

$$\sum_{t \in [T]} \mathbf{m}_t(v) \sum_{k \in K(t)} x_{tk} \ln \left(\frac{x_{tk}}{e} \right) + \sum_{i \in [d]} \alpha_i \left(\sum_{t \in [T]} \sum_{k \in K(t)} v_{itk} \cdot x_{tk} - L^* (1 + \epsilon) \right) + \sum_{t \in [T]} \beta_t \left(1 - \sum_{k \in K(t)} x_{tk} \right).$$

From the KKT conditions for the optimal solution to the convex program as a function of $x^{(\epsilon)}, \alpha^{(\epsilon)}, \beta^{(\epsilon)}$, the solution that allocates according to $x^{(\epsilon)}$ is a $(1 + \epsilon)$ -approximation to the optimal

objective L^* , and $\alpha_i^{(\epsilon)} \geq 0$ for all $i \in [d]$. Furthermore,

$$\frac{dL}{dx_{tk}} = 0 \quad \text{for all } k \in K(t), \text{ which gives } \mathbf{m}_t(v) \cdot \ln(x_{tk}^{(\epsilon)}) + \mathbf{m}_t(v) \cdot \sum_i \alpha_i^{(\epsilon)} \cdot v_{itk} = \beta_t^{(\epsilon)}.$$

For any two options $k, r \in K(t)$, we obtain:

$$\mathbf{m}_t(v) \cdot \ln(x_{tk}^{(\epsilon)}) + \sum_i \alpha_i^{(\epsilon)} \cdot v_{itk} = \mathbf{m}_t(v) \cdot \ln(x_{tr}^{(\epsilon)}) + \sum_i \alpha_i^{(\epsilon)} \cdot v_{itr}.$$

Therefore, $\ln\left(\frac{x_{tk}^{(\epsilon)}}{x_{tr}^{(\epsilon)}}\right) = \sum_i \alpha_i^{(\epsilon)} \cdot \frac{v_{itr}}{\mathbf{m}_t(v)} - \sum_i \alpha_i^{(\epsilon)} \cdot \frac{v_{itk}}{\mathbf{m}_t(v)}$. Coupled with $\sum_{k \in K(t)} x_{tk}^{(\epsilon)} = 1$, we get

$$x_{tk}^{(\epsilon)} \propto \exp\left(-\sum_i \alpha_i^{(\epsilon)} \cdot \frac{v_{itk}}{\mathbf{m}_t(v)}\right). \quad \square$$

Bounding the Learned Parameters. For $\epsilon = 0$, the learned parameters $\alpha^{(\epsilon)}$ in the previous lemma may be unbounded. However, for $\epsilon > 0$, we show that each $\alpha_i^{(\epsilon)}$ can be bounded as a function of d and ϵ . Our main tool is perturbation and sensitivity analysis, following the framework of Boyd et al. [16]. (See Appendix B for details of perturbation and sensitivity analysis.)

Lemma 4.6. *Let $x^{(\epsilon)}, \alpha^{(\epsilon)}, \beta^{(\epsilon)}$ be the optimal solution to the convex program in Figure 1, for some $\epsilon > 0$. Then, for all $i \in [d]$, it holds that $\alpha_i^{(\epsilon)} \leq \frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right)$.*

Proof. We need the following claim:

Claim 4.7. *For an n dimensional vector $x \in [0, 1]^n$ such that $\sum_{i=1}^n x_i = 1$ we have $\sum_{k=1}^n x_k \ln x_k \geq -\ln n$.*

Proof. By Jensen's inequality, we have

$$\varphi\left(\sum_i p_i \cdot y_i\right) \geq \sum_i p_i \cdot \varphi(y_i), \text{ where } p_i \geq 0, \sum_i p_i = 1, \text{ and } \varphi \text{ is concave.}$$

The desired bound now follows by setting $\varphi(x) = \ln(x)$, $p_i = x_i$ and $y_i = 1/x_i$. \square

We define a perturbed convex program based on Figure 1, where u_i corresponds to the constraint α_i .

For each $i \in [d]$, setting $u_i = -\epsilon \cdot L^*$ and $u_{i'} = 0$ for $i' \neq i$ ensures that constraint i in the perturbed problem matches the original constraint, thereby guaranteeing a feasible solution. By Lemma B.1,

$$p^*(0, 0) \geq p^*(u, v) + \alpha_i^{(\epsilon)} \cdot \epsilon \cdot L^*, \text{ which implies}$$

$$\alpha_i^{(\epsilon)} \cdot \epsilon \cdot L^* \leq p^*(0, 0) - p^*(u, v) \leq \sum_{t \in [T]} \mathbf{m}_t(v) \ln |K(t)|,$$

where the last inequality follows from Claim 4.7.

$$\sum_t \mathbf{m}_t(v) = \sum_t \mathbf{m}_t(v) \sum_k x_{tk}^* \leq \sum_t \sum_k x_{tk}^* \sum_i v_{itk} = \sum_i \sum_t \sum_k v_{itk} \cdot x_{tk}^* \leq \sum_i L^* = d \cdot L^*,$$

where the first equality holds because x^* is a feasible solution, the first inequality follows since there must be at least one nonzero coordinate and by definition of $\mathbf{m}_t(v)$, and the second inequality follows from the linear program constraint for dimension i .

By our assumption on the instance, we have

$$\ln |K(t)| \leq d \cdot \ln\left(\log_{1+\epsilon}\left(\frac{d^2}{\epsilon^2}\right)\right) \leq d \cdot \ln\left(\frac{d}{\epsilon}\right). \text{ Therefore, } \alpha_i^{(\epsilon)} \leq \frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right). \quad \square$$

Discretizing the Learned Parameters and Noise Resilience. In order to prove Lemma 4.4, we need to ensure that the learned parameters $\alpha^{(\epsilon)}$ belong to the discrete set $\text{NET}\left(\frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right), \frac{d^3}{\epsilon^3}\right)$. This may not be true of the Lagrangian multipliers derived above. We show that in discretizing the learned parameters, the competitive ratio only worsens by $(1 + O(\epsilon))$. In particular, if we replace the Lagrangian vector $\alpha^{(\epsilon)}$ with a perturbed vector $\tilde{\alpha}$, where $-\frac{\epsilon^3}{d^3} \leq \tilde{\alpha}_i - \alpha_i^{(\epsilon)} \leq \frac{\epsilon^3}{d^3}$ for all $i \in [d]$, then the assignment fractions change by at most $1 \pm 4\epsilon$. Note that this lemma is also important from a noise resilience perspective: if the learned parameters have small error, this lemma shows that the resulting allocation is still approximately optimal.

Lemma 4.8. *Let $\alpha^* \in \mathbb{R}^d$ be a vector, and let $\tilde{\alpha} \in \mathbb{R}^d$ be a perturbed vector such that $|\tilde{\alpha}_i - \alpha_i^*| \leq \frac{\epsilon^3}{d^3}$ for all $i \in [d]$. Then, the fractional assignment $x_{tk}(\tilde{\alpha})$ satisfies*

$$(1 - 4\epsilon) x_{tk}(\alpha^*) \leq x_{tk}(\tilde{\alpha}) \leq (1 + 4\epsilon) x_{tk}(\alpha^*) \quad \text{for all } k \in K(t), t \in [T].$$

Proof. We need the following claim:

Claim 4.9. *For any $\epsilon > 0$, if we are given two sets of K weights $a, a' \in \mathbb{R}_+^K$ such that $1 - \epsilon \leq \frac{a_k}{a'_k} \leq 1 + \epsilon$, then for $x, x' \in [0, 1]^K$ such that $\sum_{k=1}^K x_k = \sum_{k=1}^K x'_k = 1$ and $x_k \propto a_k$ and $x'_k \propto a'_k$, we have $1 - 4\epsilon \leq \frac{x_k}{x'_k} \leq 1 + 4\epsilon$.*

Proof. Consider the assignment fractions x_k . By definition, we have the following:

$$\begin{aligned} x_k &= \left(\sum_{r=1}^K \frac{a_r}{a_k} \right)^{-1} \leq \left(\sum_{r=1}^K \frac{a'_r}{a'_k} \cdot \frac{1 - \epsilon}{1 + \epsilon} \right)^{-1} \leq \frac{1 + \epsilon}{1 - \epsilon} \cdot \left(\sum_{r=1}^K \frac{a'_r}{a'_k} \right)^{-1} \leq (1 + 4\epsilon) \cdot x'_k. \\ x_k &= \left(\sum_{r=1}^K \frac{a_r}{a_k} \right)^{-1} \geq \left(\sum_{r=1}^K \frac{a'_r}{a'_k} \cdot \frac{1 + \epsilon}{1 - \epsilon} \right)^{-1} \geq \frac{1 - \epsilon}{1 + \epsilon} \cdot \left(\sum_{r=1}^K \frac{a'_r}{a'_k} \right)^{-1} \geq (1 - 4\epsilon) \cdot x'_k. \end{aligned} \quad \square$$

Define $a_{tk}(\alpha) = \exp\left(-\sum_i \alpha_i \cdot \frac{v_{itk}}{\mathbf{m}_t(v)}\right)$. Then, we have

$$\frac{a_{tk}(\alpha^*)}{a_{tk}(\tilde{\alpha})} = \exp\left(-\sum_i (\alpha_i - \tilde{\alpha}_i) \cdot \frac{v_{itk}}{\mathbf{m}_t(v)}\right).$$

By the assumption that the instance is balanced, we have $\frac{v_{itk}}{\mathbf{m}_t(v)} \leq \frac{d^2}{\epsilon^2}$, which implies

$$1 - 2\epsilon \leq \exp(-\epsilon) \leq \frac{a_{tk}(\alpha^*)}{a_{tk}(\tilde{\alpha})} \leq \exp(\epsilon) \leq 1 + 2\epsilon.$$

Thus, by Claim 4.9, the assignment fractions using $\tilde{\alpha}$ differ from those using α by at most $(1 \pm 4\epsilon)$. \square

Finally, we now put all the pieces together to establish Lemma 4.4:

Proof of Lemma 4.4. Fix a balanced instance $I(v, K)$ and $\epsilon > 0$. By Lemma 4.6, there exists a parameter vector $\alpha^{(\epsilon)}$ such that $\alpha_i^{(\epsilon)} \in \left[0, \frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right)\right]$ for all $i \in [d]$. Therefore, there exists a vector $\tilde{\alpha} \in \text{NET}\left(\frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right), \frac{d^3}{\epsilon^3}\right)$ such that $|\tilde{\alpha}_i - \alpha_i^{(\epsilon)}| \leq \frac{\epsilon^3}{d^3}$ for all $i \in [d]$. By Lemma 4.8, the exponential assignment rule with $\tilde{\alpha}$ achieves a $(1 + O(\epsilon))$ -approximation. \square

5 Applications

Our framework applies broadly to a variety of online allocation problems. We briefly highlight two representative applications: online routing and Nash Social Welfare maximization.

Online Routing. In online routing (e.g., [3]), requests arrive over time and must be assigned to paths in a network to minimize congestion or delay. Each option (path) contributes load to

network edges, and the goal is to minimize the maximum edge utilization, which corresponds to the **MinMax** objective. Our approach uses a vector of learned edge-weight parameters to guide allocation decisions, combining strong worst-case guarantees with improved performance when input patterns are predictable.

Formally, we are given as input a directed graph $G = (V, E)$. At each time step $t \in [T]$, a flow request of amount r_t between two vertices s_t and t_t , along with a set of candidate paths $\mathcal{P}_t = \{P_{t,1}, P_{t,2}, \dots\}$, is revealed. Here, each $P_{t,k}$ is a path from s_t to t_t in the graph. The algorithm must assign values $q_{t,k}$ to each path such that $\sum_k q_{t,k} = r_t$. The load on an edge $e \in E$ is given by $\ell_e(q) = \sum_t \sum_{P_{t,k}|e \in P_{t,k}} q_{t,k}$, and the objective is to minimize the maximum congestion on any edge, i.e., $\min \max_{e \in E} \{\ell_e(q) \mid \sum_k q_{t,k} = r_t \text{ for all } t \in [T]\}$.

We note that online routing is a special case of the online allocation problem, where $A_t \subseteq \left\{ z_t \cdot \vec{f} \mid \vec{f} \in \{0,1\}^d, z_t \in \mathbb{R}_+ \right\}$, and the objective is to minimize $f(v) = \max_{i \in [d]} v_i$.

Our main result for online routing follows as a corollary of Theorem 4.1.

Corollary 5.1. *Given an instance of online routing and any $\epsilon > 0$, there exists a set of parameters $\alpha \in \text{NET}\left(\frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right), \frac{d^3}{\epsilon^3}\right)$ such that the fractional solution $x_{t,k} \propto \exp(-\sum_i \alpha_i \cdot f_{itk})$, for $t \in [T]$, approximates the optimal objective within a factor of $(1 + O(\epsilon))$.*

Online Nash Social Welfare. As mentioned, our scheme also applies to maximization problems where, at each step, a divisible resource has to be distributed among a set of d agents. The objective is to design an online algorithm that balances fairness and efficiency. At the start of each step t , the algorithm observes the value $v_{i,t}$ of each agent i for that resource, and then irrevocably determines the allocation without knowledge of future values. If agent i is allocated a fraction $x_{i,t}$, their utility increases by $x_{i,t}v_{i,t}$. The total utility of agent i is then given by $u_i(\mathbf{x}) = \sum_t v_{i,t}x_{i,t}$.

The *Nash Social Welfare* (NSW) objective is known to provide a natural balance between fairness and efficiency. It is defined as the geometric mean of the agents' utilities: $\text{NSW}(\mathbf{x}) = (\prod_i u_i(\mathbf{x}))^{1/d}$. By applying our method for maximizing well-behaved objectives, we obtain the following corollary:

Corollary 5.2. *Given an instance of online NSW maximization and any $\epsilon > 0$, there exists a set of parameters $\alpha \in \text{NET}\left(\frac{d^2}{\epsilon} \cdot \ln\left(\frac{d}{\epsilon}\right), \frac{d^3}{\epsilon^3}\right)$ such that the fractional solution $x_{t,k} \propto \exp(\sum_i \alpha_i \cdot v_{i,t,k})$, for $t \in [T]$, approximates the optimal objective within a factor of $(1 - O(\epsilon))$.*

Note that using our results, we may further generalize this to options with vector utilities, i.e., where an option can add (possibly different amounts of) value to multiple agents simultaneously.

6 Closing Remarks

In this paper, we gave a general technique for designing online algorithms with predictions that applies to the entire spectrum of problems that can be modeled via the online covering framework. This has value in two respects. First, it gave the first learning-augmented results for important problems like online routing and scheduling that goes beyond simple assignment, and were beyond the scope of previous techniques. Second, and perhaps more importantly, it opens the door to the design of even more general-purpose methods for the design of online algorithms with predictions. For instance, can we give a technique for learning-augmented algorithms that applies to any covering LP where each online step reveals a new constraint in the LP? Techniques such as the online primal dual method that apply to this class of LPs have been hugely influential in the classical online algorithms literature (without predictions), which makes it a tempting proposition to explore similarly powerful tools in the learning-augmented setting.

Acknowledgments and Disclosure of Funding

Ilan Reuven Cohen's research was supported by the Israel Science Foundation grant No. 1737/21. Debmalya Panigrahi's research was supported in part by NSF grants CCF-1955703 and CCF-2329230.

References

- [1] Shipra Agrawal and Nikhil R Devanur. Fast algorithms for online stochastic convex programming. In *Proceedings of the twenty-sixth annual ACM-SIAM symposium on Discrete algorithms*, pages 1405–1424. SIAM, 2014.
- [2] Antonios Antoniadis, Themis Gouleakis, Pieter Kleer, and Pavel Kolev. Secretary and online matching problems with machine learned advice. In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020*, 2020.
- [3] James Aspnes, Yossi Azar, Amos Fiat, Serge Plotkin, and Orli Waarts. On-line routing of virtual circuits with applications to load balancing and machine scheduling. *Journal of the ACM (JACM)*, 44(3):486–504, 1997.
- [4] Baruch Awerbuch and Yossi Azar. Competitive multicast routing. *Wirel. Networks*, 1(1):107–114, 1995.
- [5] Baruch Awerbuch, Yossi Azar, Edward F. Grove, Ming-Yang Kao, P. Krishnan, and Jeffrey Scott Vitter. Load balancing in the l_p norm. In *36th Annual Symposium on Foundations of Computer Science, Milwaukee, Wisconsin, USA, 23-25 October 1995*, pages 383–391. IEEE Computer Society, 1995.
- [6] Baruch Awerbuch, Yossi Azar, and Serge A. Plotkin. Throughput-competitive on-line routing. In *34th Annual Symposium on Foundations of Computer Science, Palo Alto, California, USA, 3-5 November 1993*, pages 32–40. IEEE Computer Society, 1993.
- [7] Baruch Awerbuch, Yossi Azar, Serge A. Plotkin, and Orli Waarts. Competitive routing of virtual circuits with unknown duration. In Daniel Dominic Sleator, editor, *Proceedings of the Fifth Annual ACM-SIAM Symposium on Discrete Algorithms. 23-25 January 1994, Arlington, Virginia, USA*, pages 321–327. ACM/SIAM, 1994.
- [8] Yossi Azar, Stefano Leonardi, and Noam Touitou. Flow time scheduling with uncertain processing time. In *STOC ’21: 53rd Annual ACM SIGACT Symposium on Theory of Computing*, pages 1070–1080. ACM, 2021.
- [9] Yossi Azar, Stefano Leonardi, and Noam Touitou. Distortion-oblivious algorithms for minimizing flow time. In *Proceedings of the 2022 ACM-SIAM Symposium on Discrete Algorithms, SODA 2022*, pages 252–274. SIAM, 2022.
- [10] Yossi Azar, Joseph Naor, and Raphael Rom. The competitiveness of on-line assignments. *J. Algorithms*, 18(2):221–237, 1995.
- [11] Étienne Bamas, Andreas Maggiori, Lars Rohwedder, and Ola Svensson. Learning augmented energy minimization via speed scaling. In *Advances in Neural Information Processing Systems 33, NeurIPS 2020*, 2020.
- [12] Etienne Bamas, Andreas Maggiori, and Ola Svensson. The primal-dual method for learning augmented algorithms. *Advances in Neural Information Processing Systems*, 33:20083–20094, 2020.
- [13] Siddhartha Banerjee, Vasilis Gkatzelis, Artur Gorokh, and Billy Jin. Online nash social welfare maximization with predictions. In *Proceedings of the 2022 ACM-SIAM Symposium on Discrete Algorithms, SODA 2022*, pages 1–19. SIAM, 2022.
- [14] Siddhartha Banerjee, Vasilis Gkatzelis, Artur Gorokh, and Billy Jin. Online nash social welfare maximization with predictions. In *Proceedings of the 2022 Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*, pages 1–19. SIAM, 2022.
- [15] Siddharth Barman, Arindam Khan, and Arnab Maiti. Universal and tight online algorithms for generalized-mean welfare. In *Thirty-Sixth AAAI Conference on Artificial Intelligence*, pages 4793–4800. AAAI Press, 2022.

- [16] Stephen Boyd and Lieven Vandenberghe. *Convex optimization*. Cambridge university press, 2004.
- [17] Justin Y. Chen and Piotr Indyk. Online bipartite matching with predicted degrees. *CoRR*, 2021.
- [18] Ilan Reuven Cohen and Debmalya Panigrahi. A general framework for learning-augmented online allocation. In Kousha Etessami, Uriel Feige, and Gabriele Puppis, editors, *50th International Colloquium on Automata, Languages, and Programming, ICALP 2023, July 10-14, 2023, Paderborn, Germany*, volume 261 of *LIPICS*, pages 43:1–43:21. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2023.
- [19] Elena Grigorescu, Young-San Lin, Sandeep Silwal, Maoyuan Song, and Samson Zhou. Learning-augmented algorithms for online linear and semidefinite programming. *Advances in Neural Information Processing Systems*, 35:38643–38654, 2022.
- [20] Elena Grigorescu, Young-San Lin, and Maoyuan Song. Learning-augmented algorithms for online concave packing and convex covering problems. *arXiv preprint arXiv:2411.08332*, 2024.
- [21] Sungjin Im, Ravi Kumar, Mahshid Montazer Qaem, and Manish Purohit. Non-clairvoyant scheduling with predictions. In *SPAA '21: 33rd ACM Symposium on Parallelism in Algorithms and Architectures, Virtual Event, USA, 6-8 July, 2021*, pages 285–294. ACM, 2021.
- [22] Ravi Kumar, Manish Purohit, Aaron Schild, Zoya Svitkina, and Erik Vee. Semi-online bipartite matching. In *10th Innovations in Theoretical Computer Science Conference, ITCS 2019*, volume 124 of *LIPICS*, pages 50:1–50:20. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2019.
- [23] Silvio Lattanzi, Thomas Lavastida, Benjamin Moseley, and Sergei Vassilvitskii. Online scheduling via learned weights. In *Proceedings of the 2020 ACM-SIAM Symposium on Discrete Algorithms, SODA 2020*, pages 1859–1877. SIAM, 2020.
- [24] Thomas Lavastida, Benjamin Moseley, R. Ravi, and Chenyang Xu. Learnable and instance-robust predictions for online matching, flows and load balancing. In *29th Annual European Symposium on Algorithms, ESA 2021*, volume 204 of *LIPICS*, pages 59:1–59:17, 2021.
- [25] Thomas Lavastida, Benjamin Moseley, R. Ravi, and Chenyang Xu. Using predicted weights for ad delivery. In *Applied and Computational Discrete Algorithms, ACDA 2021*, 2021.
- [26] Shi Li and Jiayi Xian. Online unrelated machine load balancing with predictions revisited. In *Proceedings of the 38th International Conference on Machine Learning, ICML 2021*, 2021.
- [27] Thodoris Lykouris and Sergei Vassilvitskii. Competitive caching with machine learned advice. *J. ACM*, 68(4):24:1–24:25, 2021.
- [28] Mohammad Mahdian, Hamid Nazerzadeh, and Amin Saberi. Online optimization with uncertain information. *ACM Trans. Algorithms*, 8(1):2:1–2:29, 2012.
- [29] Aranyak Mehta, Amin Saberi, Umesh Vazirani, and Vijay Vazirani. Adwords and generalized online matching. *Journal of the ACM (JACM)*, 54(5):22–es, 2007.
- [30] Michael Mitzenmacher. Scheduling with predictions and the price of misprediction. In *11th Innovations in Theoretical Computer Science Conference, ITCS 2020*, volume 151 of *LIPICS*, pages 14:1–14:18. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2020.
- [31] Michael Mitzenmacher and Sergei Vassilvitskii. Algorithms with predictions. In *Beyond the Worst-Case Analysis of Algorithms*, pages 646–662. Cambridge University Press, 2020.
- [32] Michael Mitzenmacher and Sergei Vassilvitskii. Algorithms with predictions. *Commun. ACM*, 65(7):33–35, 2022.
- [33] Manish Purohit, Zoya Svitkina, and Ravi Kumar. Improving online algorithms via ML predictions. In *Advances in Neural Information Processing Systems 31, NeurIPS 2018*, 2018.

- [34] Max Springer, MohammadTaghi Hajiaghayi, Debmalya Panigrahi, and Mohammad Reza Khani. Online algorithms for the santa claus problem. In Sanmi Koyejo, S. Mohamed, A. Agarwal, Danielle Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*, 2022.
- [35] Nguyen Kim Thang and Christoph Durr. Online primal-dual algorithms with predictions for packing problems. *arXiv preprint arXiv:2110.00391*, 2021.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: The paper develops a learning-augmented scheme for online allocation and thoroughly proves its guarantees and applicability.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Yes, the paper discusses the limitations by outlining the modeling choices and assumptions under which the results hold. It clarifies the settings where the approach is applicable, such as the well-behaved objective functions and distributional assumptions for learnability.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Yes. The paper clearly states all necessary assumptions for each theoretical result and provides complete and correct proofs throughout.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [NA]

Justification: The paper does not include experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [NA]

Justification: The paper does not include experiments requiring code.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [NA]

Justification: The paper does not include experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [NA]

Justification: The paper does not include experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).

- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [NA]

Justification: The paper does not include experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Yes.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: There is no societal impact of the work performed.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: The paper poses no such risks.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [NA]

Justification: The paper does not use existing assets.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: The paper does not release new assets.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: The paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: the paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification: The core method development in this research does not involve LLMs.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.

A Preprocessing for the MinMax Objective

We begin by introducing a preprocessing step for the **MinMax** objective. Given an instance $I(v, K)$, let $\text{MinMax}(v, K)$ denote the optimal value of the corresponding **MinMax** objective. Specifically, we construct an instance $I(\tilde{v}, \tilde{K})$ that ϵ -approximates a given instance $I(v, K)$ as follows:

Definition A.1. An instance $I(\tilde{v}, \tilde{K})$ is said to ϵ -approximate an instance $I(v, K)$ if:

- $\tilde{K}(j) \subseteq K(j)$ for all $j \in N$,
- $\text{MinMax}(\tilde{v}, \tilde{K}) \leq \text{MinMax}(v, K) \cdot (1 + \epsilon)$,
- For any feasible allocation \tilde{x} for $I(\tilde{v}, \tilde{K})$, if the load in any dimension is at most L with respect to \tilde{v} , then the corresponding load with respect to v is at most $L \cdot (1 + \epsilon)$.

(Note that \tilde{x} remains a feasible allocation for $I(v, K)$ since $\tilde{K}(j) \subseteq K(j)$.)

Using this definition, we prove that in order to achieve a $(1 + O(\epsilon))$ to the instance $I(v, K)$ it is sufficient to achieve $(1 + \epsilon)$ approximation to $I(\tilde{v}, \tilde{K})$.

Corollary A.2. Let $I(v, K)$ be a given instance. Suppose the instance is transformed into an ϵ -approximate instance $I(\tilde{v}, \tilde{K})$. If a $(1 + \epsilon)$ -approximate allocation is computed for the transformed instance $I(\tilde{v}, \tilde{K})$, then this allocation guarantees a $(1 + O(\epsilon))$ -approximation for the original instance I .

Proof. Let x be a feasible solution for $I(\tilde{v}, \tilde{K})$ such that

$$\sum_t \sum_{k \in \tilde{K}(t)} x_{tk} \cdot \tilde{v}_{itk} \leq \text{MinMax}(\tilde{v}, \tilde{K}) \cdot (1 + \epsilon).$$

Using Definition A.1, we obtain:

$$\begin{aligned} \sum_t \sum_{k \in \tilde{K}(t)} x_{tk} \cdot v_{itk} &\leq \text{MinMax}(\tilde{v}, \tilde{K}) \cdot (1 + \epsilon)^2 \\ &\leq \text{MinMax}(v, K) \cdot (1 + \epsilon)^3. \end{aligned}$$

□

We now show that each of the three steps defined in the transformation produces an ϵ -approximate instance with respect to the original instance.

Claim A.3 (Step (1) of the preprocessing). *Given an instance $I(v, K)$ let \hat{v} such that*

$$\hat{v}_{itk} = \begin{cases} 0 & \text{if } \frac{v_{itk}}{\max_{i'} v_{i'tk}} < \frac{\epsilon}{d}, \\ v_{itk} & \text{otherwise.} \end{cases}$$

the the instance $I(\hat{v}, K)$ ϵ -approximate the instance $I(v, K)$.

Proof. By definition, $\text{MinMax}(\tilde{v}, \tilde{K}) \leq \text{MinMax}(v, K)$. Next, given a solution x such that for all $i \in [d]$,

$$\sum_{tk} x_{tk} \cdot \tilde{v}_{itk} \leq L,$$

we prove that

$$\sum_t \sum_{k \in \tilde{K}(t)} x_{tk} \cdot v_{itk} \leq L^* \cdot (1 + \epsilon).$$

For a fixed coordinate $i \in [d]$, define

$$R_i(t) = \{k \in K(t) : v_{itk} > \tilde{v}_{itk}\},$$

which represents the set of options for item t where the load on coordinate i is modified by the transformation. For an item t and an option $k \in K(t)$, let

$$i_{tk}^{\max} \in \arg \max_{i' \in [d]} v_{i'tk}$$

denote the coordinate that experiences the maximum load for option k of item t .

Observe that if $k \in R_i(t)$, then

$$v_{itk} \leq \frac{\epsilon}{d} \cdot v_{i_{tk}^{\max}, t, k}.$$

Thus, we have:

$$\begin{aligned} & \sum_{tk} x_{tk} \cdot v_{itk} \\ &= \sum_{tk} x_{tk} \cdot \tilde{v}_{itk} + \sum_t \sum_{k \in R_i(t)} x_{tk} \cdot v_{itk} \\ &\leq L + \sum_t \sum_{k \in R_i(t)} x_{tk} \cdot v_{itk} \\ &= L + \sum_{i' \in d} \sum_t \sum_{k \in R_i(t) | i' = i_{tk}^{\max}} x_{tk} \cdot v_{itk} \\ &\leq L + \sum_{i' \in d} \sum_t \sum_{k \in R_i(t) | i' = i_{tk}^{\max}} x_{tk} \cdot \frac{\epsilon}{d} \cdot \tilde{v}_{i'tk} \\ &= L + \sum_{i' \in d} \sum_t \sum_{k \in R_i(t) | i' = i_{tk}^{\max}} \frac{\epsilon}{d} \cdot x_{tk} \cdot \tilde{v}_{i'tk} \\ &\leq L + \frac{\epsilon}{d} \cdot \sum_{i' \in d} \sum_{tk} x_{tk} \cdot \tilde{v}_{i'tk} \\ &\leq L + \sum_{i' \in [d]} L \cdot \frac{\epsilon}{d} \\ &= L \cdot (1 + \epsilon), \end{aligned}$$

where the first and final inequalities follow from the definition of x , and the second inequality follows from the transformation definition. \square

The

Claim A.4 (Step (2) of the Preprocessing). *Given an instance $I(v, K)$, let $I(\tilde{v}, K')$ be the instance obtained by removing option k' from item t whenever there exists another option $k \in K(t)$ such that*

$$v_{t,k'}^{\max} > \frac{d}{\epsilon} \cdot v_{t,k}^{\max},$$

where $v_{t,k}^{\max} = \max_{i \in [d]} v_{i,t,k}$.

Then $I(v, K')$ is an ϵ -approximation of the instance $I(v, K)$.

Proof. Clearly, an assignment with a load of at most L for all $i \in [d]$ in $I(\tilde{v}, \tilde{K})$ results in a load of at most L in $I(v, K)$, since the retained options have identical loads. We now show that if an allocation x achieves a makespan of L in $I(v, K)$, then there exists a transformed allocation \hat{x} that achieves a makespan of at most $L \cdot (1 + \epsilon)$ in $I(\tilde{v}, \tilde{K})$.

We explicitly construct such an assignment \hat{x} and show that it satisfies the required bound. Define

$$k_t^{\min} \in \arg \min_k v_{t,k}^{\max}$$

as the option with the smallest maximum load for item t , breaking ties by selecting the smallest index. Let

$$i_{tk}^{\max} = \arg \max_i v_{i,t,k}$$

denote the coordinate where option k of item t imposes the maximum load.

Next, define the set

$$R(t) = \{k \in K(t) \mid v_{t,k}^{\max} > \frac{d}{\epsilon} \cdot v_{t,k_t^{\min}}^{\max}\},$$

which identifies options with disproportionately high loads relative to the smallest maximum load. We then construct the transformed allocation \hat{x} for each item $t \in [n]$ as follows:

$$\hat{x}_{tk} = \begin{cases} x_{tk} + \sum_{k' \in R(t)} x_{tk'} & \text{if } k = k_t^{\min}, \\ 0 & \text{if } k \in R(t), \\ x_{tk} & \text{otherwise.} \end{cases}$$

This transformation ensures that options in $R(t)$ are reallocated to the option with the smallest maximum load, maintaining feasibility while ensuring that the makespan increases by at most a factor of $(1 + \epsilon)$. Specifically, we have:

$$\begin{aligned} & \sum_{tk} \hat{x}_{tk} \cdot v_{itk} \\ &= \sum_{tk} x_{tk} \cdot v_{itk} + \sum_t \sum_{k' \in R(t)} x_{tk'} \cdot v_{i,t,k_t^{\min}} \\ &\leq L + \sum_t \sum_{k' \in R(t)} x_{tk'} \cdot \frac{\epsilon}{d} \cdot v_{i_{tk'}^{\max}, t, k'} \\ &= L + \frac{\epsilon}{d} \cdot \sum_{i', t} \sum_{k' \in R(t)} \sum_{i'=i_{tk'}^{\max}} x_{tk'} \cdot v_{i', t, k'} \\ &\leq L + \frac{\epsilon}{d} \cdot \sum_{i'} \sum_{t, k} x_{tk} \cdot v_{i', t, k} \\ &\leq L + \sum_{i' \in [d]} L \cdot \frac{\epsilon}{d} \\ &= L \cdot (1 + \epsilon), \end{aligned}$$

where the first and final inequalities follow from the definition of x , and the second inequality follows from the fact that for $k' \in R(t)$,

$$v_{i,t,k_t^{\min}} \leq v_{t,k_t^{\min}}^{\max} < \frac{\epsilon}{d} \cdot v_{t,k'}^{\max} = \frac{\epsilon}{d} \cdot v_{i_{tk'}^{\max}, t, k'}.$$

□

Claim A.5 (Step (3) of the Preprocessing). *Given an instance $I(v, K)$*

if $v_{itk} = 0$, then set $\tilde{v}_{itk} = 0$; otherwise, set

$$\tilde{v}_{itk} = \mathbf{m}_t(v) \cdot (1 + \epsilon)^{\lfloor \log_{1+\epsilon} \left(\frac{\tilde{v}_{itk}}{\mathbf{m}_t(v)} \right) \rfloor}.$$

Then $I(\tilde{v}, K)$ is an ϵ -approximation of the instance $I(v, K)$.

Proof. By definition,

$$\tilde{v}_{itk} \leq v_{itk} \leq (1 + \epsilon) \cdot \tilde{v}_{itk}.$$

Therefore, $I(\tilde{v}, K)$ is an ϵ -approximation of the instance $I(v, K)$.

□

By applying Claim A.3, Claim A.4, Claim A.5, and Corollary A.2 we proved the Lemma 4.3.

B Perturbation and Sensitivity Analysis

One of our main tools for bounding learned parameters is perturbation and sensitivity analysis, following the framework of Boyd et al. [16].

B.1 Perturbed Convex Programs

Consider a convex program represented by a tuple (f, h) , where $f_i : \mathbb{R}^d \rightarrow \mathbb{R}$ for $i \in [0, \tilde{m}]$ and $h_i : \mathbb{R}^d \rightarrow \mathbb{R}$ for $i \in [1, \tilde{n}]$, formulated as:

$$\begin{aligned} & \text{minimize } f_0(x) \\ & \text{subject to } f_i(x) \leq 0, \quad i \in [1, \tilde{m}] \\ & \quad h_i(x) = 0, \quad i \in [1, \tilde{n}]. \end{aligned}$$

We define the perturbed version of this problem as follows. Let $u \in \mathbb{R}^{\tilde{m}}$ and $v \in \mathbb{R}^{\tilde{n}}$. Define perturbed constraints:

$$f'_i(x) = f_i(x) - u_i, \quad h'_i(x) = h_i(x) - v_i.$$

The perturbed problem becomes:

$$\begin{aligned} & \text{minimize } f_0(x) \\ & \text{subject to } f_i(x) \leq u_i, \quad i \in [1, \tilde{m}] \\ & \quad h_i(x) = v_i, \quad i \in [1, \tilde{n}]. \end{aligned}$$

This coincides with the original problem when $u = 0$ and $v = 0$. Positive u_i relaxes the i th inequality; negative u_i tightens it. The vector v perturbs the right-hand sides of the equality constraints.

Let $p^*(u, v)$ denote the optimal value of the perturbed problem. If the problem is infeasible, we define $p^*(u, v) = \infty$. When the original problem is convex, $p^*(u, v)$ is convex in both u and v .

B.2 A Global Inequality via Duality

Assume strong duality holds and that the dual optimum is attained (which is the case under Slater's condition). Let (λ^*, ν^*) denote an optimal dual solution to the original problem. Then the following global bound holds.

Lemma B.1 (Perturbation Inequality, [16]). *For all $u \in \mathbb{R}^{\tilde{m}}$ and $v \in \mathbb{R}^{\tilde{n}}$,*

$$p^*(0, 0) \geq p^*(u, v) - \lambda^{*T} u - \nu^{*T} v.$$

B.3 Proof Sketch

To derive this inequality, consider any feasible solution x to the perturbed problem, i.e.,

$$f_i(x) \leq u_i, \quad \forall i \in [1, \tilde{m}], \quad h_i(x) = v_i, \quad \forall i \in [1, \tilde{n}].$$

By strong duality:

$$\begin{aligned} p^*(0, 0) &= g(\lambda^*, \nu^*) \\ &\leq f_0(x) + \sum_{i=1}^{\tilde{m}} \lambda_i^* f_i(x) + \sum_{i=1}^{\tilde{n}} \nu_i^* h_i(x) \\ &\leq f_0(x) + \lambda^{*T} u + \nu^{*T} v, \end{aligned}$$

where the last inequality follows from $f_i(x) \leq u_i$, $h_i(x) = v_i$, and $\lambda^* \geq 0$.

Rearranging, we conclude:

$$f_0(x) \geq p^*(0, 0) - \lambda^{*T} u - \nu^{*T} v.$$

This inequality quantifies how much the objective value may change in response to perturbations of the constraints, depending linearly on the dual multipliers.

C Learnability of the Prediction Vectors

We consider the learning model introduced by [24], and show that under this model, the dual vector α can be learned efficiently from sampled instances. Specifically, we consider the following model: the t th step (i.e., the values of $v_t = (v_{i,t,k} : i \in [d], k \in K(t))$) is independently sampled from a (discrete) distribution \mathcal{D}_j .

We set up the online allocation problem for the **MinMax** objective; the setup for the **MaxMin** objective is very similar and is omitted for brevity. Let $L = \mathbb{E}_{P \sim \mathcal{D}}[\text{MinMax}(P)]$ be the expected value of the **MinMax** objective in the optimal solution for an instance P drawn from \mathcal{D} .

Morally, we would like to say that we can obtain the vector α that gives a nearly optimal solution (in expectation) using vector allocation (i.e., a **MinMax** objective of $(1 + \epsilon) \cdot L$ in expectation for some error parameter ϵ) using a bounded (as a function of ϵ) number of samples. Similar to [24], we need the following assumption:

Small Items Assumption: Conceptually, this assumption states that each individual item has a small utility compared to the overall utility of any agent in an optimal solution. Precisely, we need $v_{itk} \leq \frac{L}{\zeta}$ for every $i \in [d], t \in [T]$, and $k \in K(t)$ for some value $\zeta = \Theta\left(\frac{\log d}{\epsilon^2}\right)$.

We show the following PAC learning theorem for the **MinMax** objective:

Theorem C.1. *Fix an $\epsilon > 0$ for which the small items assumption holds. Then, there is an (learning) algorithm that samples $O\left(\frac{d}{\log d} \cdot \log \frac{d}{\epsilon}\right)$ independent instances from \mathcal{D} and outputs (with high probability) a prediction vector α such that using α in the allocation scheme gives a **MinMax** objective of at least $(1 + O(\epsilon)) \cdot L$ in expectation over instances $P \sim \mathcal{D}$.*

Proof Sketch. Recall that in PAC theory, the number of samples needed to learn a function from a family of N functions is about $O(\log N)$. Indeed, restricting α to be in the class **NET**(K, S) serves this role of limiting the hypothesis class to a finite, bounded set since $|\text{NET}(K, S)| = (K \cdot S)^d$ where $S = K = O(\text{poly}(d, \epsilon))$. Using standard PAC theory, this implies that using about $O(d \log K) = O(d \cdot \log \frac{d}{\epsilon})$ samples, we can learn the “best” vector in **NET**(K, S). Our main technical work is to show that this “best” vector produces an approximately optimal solution when used. \square

C.1 Details of PAC Learning

Given an instance P , Let $\ell_i(P, \alpha)$ the value of the i th dimension of $\mathbf{v}_{\text{tot}}^f$ after applying our scheme on the instance P with the parameters α . Let $K = \frac{d^2}{\epsilon}$ and $S = \frac{d^3}{\epsilon^3}$.

Let us consider a combination of all instances in the support of the distribution \mathcal{D} . For L processing matrices $P^{(1)}, P^{(2)}, \dots, P^{(L)}$. We define $P^{\text{all}} = \bigoplus_{r=1}^L P^{(r)}$ to be the instance defined by the $n \cdot L$ items. For every $\ell \in [L]$ and $t \in [T]$, we have a step $t^{(\ell)}$ with values $v_t^{(\ell)}$.

The following observation is immediate (subadditivity):

Observation C.2. $\text{MinMax}(P^{\text{all}}) \leq \sum_{r=1}^L \text{MinMax}(P^{(r)})$.

Using this observation, we can prove the following lemma, by considering the combination of all instances in \mathcal{D} , scaled by their respective probabilities.

Lemma C.3. *There exists $\alpha \in \text{NET}(K, S)$, such that for every $i \in [d]$, we have*

$$\mathbb{E}_{P \sim \mathcal{D}}[\ell_i(\alpha)] \leq (1 + \epsilon) \cdot L^*.$$

Proof. Consider the instance $\mathbb{P} = \bigoplus \Pr_D[P] \cdot P$ where $\Pr_D[P]$ is the probability mass of P in \mathcal{D} , and $\Pr_D[P] \cdot P$ is the matrix P multiplied by $\Pr_D[P]$. By Observation C.2, we have

$$\begin{aligned} \text{MinMax}(\mathbb{P}) &\leq \sum_P \Pr_D[P] \text{MinMax}(P) \\ &= \mathbb{E}_{P \sim \mathcal{D}}[\text{MinMax}(P)] = L. \end{aligned}$$

We can apply Lemma 4.4 to the combined instance to show there exists $\alpha^* \in \text{NET}(K, S)$ such that for every $i \in [d]$, we have,

$$\sum_{tk} x_{t,k}(\mathbb{P}, \alpha^*) \cdot p_{itk} \leq (1 + \epsilon) \text{MinMax}(\mathbb{P}) \leq (1 + \epsilon) \cdot L^*$$

where t indexes over all steps in \mathbb{P} . Notice that $x_{t,k}(\mathbb{P}, \alpha^*)$ depends on the utility vector for step t , which is part of the instance $P \in \mathcal{D}$ that t belongs to. Therefore, the left side of the above inequality is exactly

$$\begin{aligned} & \sum_P \sum_{t,k} x_{t,k}(P, \alpha^*) \cdot \Pr_{\mathcal{D}}[P] \cdot p_{i,t,k} \\ &= \mathbb{E}_{P \sim \mathcal{D}} \sum_{t,k} x_{t,k}(P, \alpha^*) p_{i,t,k} \\ &= \mathbb{E}_{P \sim \mathcal{D}} \sum_{t \in [T]} \ell_i(P, \alpha^*), \end{aligned}$$

as required. \square

For any real numbers A, B, ϵ, C , we use $A \approx_{\epsilon,C} B$ to denote $|A - B| \leq \epsilon \cdot \max(B, C)$. The next lemma appears in [26]:

Lemma C.4 (Lemma D.6 in [26]). *For any $\alpha \in \text{NET}(K, S)$, with high probability over $P \sim \mathcal{D}$, we have*

$$\forall i \in [d] : \ell_i(P, \alpha) \approx_{\epsilon, L^*} \mathbb{E}_{P' \sim \mathcal{D}} \ell_i(P', \alpha).$$

The learning algorithm. We sample $H = O\left(\frac{d}{\log d} \log \frac{d}{\epsilon}\right)$ instances P_1, P_2, \dots, P_H independently and randomly form \mathcal{D} . We output $\tilde{\alpha} \in \text{NET}(K, S)$ that maximizes $\min_{i \in [d]} \frac{1}{H} \sum_{h=1}^H \ell_i(P_h, \tilde{\alpha})$.

The next lemma also appears in [26]:

Lemma C.5 (Lemma D.7 in [26]). *With probability at least $1 - \frac{1}{(K \cdot S)^m}$, for every $\alpha \in \text{NET}(K, S)$ and for every $i \in [d]$, we have*

$$\frac{1}{H} \sum_{h=1}^H \ell_i(P_h, \alpha) \approx_{\epsilon, L^*} \mathbb{E}_{P \sim \mathcal{D}} \ell_i(P, \alpha).$$

Now assume the event in Lemma C.5 happens. Then by Lemma C.3, there exists some $\alpha \in \text{NET}(K, S)$ such that

$$\min_{i \in [m]} \frac{1}{H} \sum_{h=1}^H \ell_i(P_h, \alpha) \leq (1 + \epsilon)^2 \cdot L^*.$$

In particular, since $\tilde{\alpha}$ maximizes $\min_{i \in [m]} \frac{1}{H} \sum_{h=1}^H \ell_i(P_h, \tilde{\alpha})$ for $\tilde{\alpha} \in \text{NET}(m, \epsilon)$, we can conclude that

$$\min_{i \in [m]} \frac{1}{H} \sum_{h=1}^H \ell_i(P_h, \tilde{\alpha}) \leq (1 + \epsilon)^2 \cdot L^*.$$

Applying Lemma C.5 again, we get

$$\min_{i \in [m]} \mathbb{E}_{P \sim \mathcal{D}} \ell_i(P, \tilde{\alpha}) \leq (1 + \epsilon)^3 \cdot L^*.$$

We now apply Lemma C.4 to $\tilde{\alpha}$. We have that with high probability over $P \sim \mathcal{D}$, for every $i \in [m]$ the following holds:

$$\begin{aligned} \ell_i(P, \tilde{\alpha}) &\leq \\ &\mathbb{E}_{P' \sim \mathcal{D}} \ell_i(P', \tilde{\alpha}) + \epsilon \cdot \max\{L^*, \mathbb{E}_{P' \sim \mathcal{D}} \ell_i(P', \tilde{\alpha})\} \leq \\ &(1 + \epsilon)^4 \cdot L^*. \end{aligned}$$

Therefore, $\text{MinMax}(P, \tilde{\alpha}) \leq (1 + \Omega(\epsilon)) \cdot L^*$. This completes the proof of Theorem C.1.

D Robustness–Consistency Tradeoff

In this section, we show that our learning-augmented scheme can be modified to balance consistency and robustness, achieving near-optimal performance when predictions are accurate while retaining strong worst-case guarantees when they are not.

Recall that an algorithm is said to be γ -consistent and δ -robust if it achieves a γ -approximation under accurate predictions (consistency), and a δ -approximation in the worst case when predictions are unreliable (robustness).

For allocation problems with minimization objectives, the worst-case approximation ratio without any predictions is $O(\log d)$, and for maximization objectives it is $O(d)$, where d is the number of agents.

We show that our learning-augmented scheme can be modified to satisfy this robustness–consistency tradeoff.

Modified Algorithm for Minimization Objectives. Let $\alpha \in \mathbb{R}^d$ be the predicted parameter vector. The algorithm operates in two phases:

1. **Prediction Phase:** At each time step, use the exponential assignment scheme with parameters α .
2. **Fallback Phase:** Monitor the cumulative objective value. If it exceeds the optimal value by a factor larger than $O(\log d)$, the algorithm switches to a standard worst-case online algorithm.

Let η be the approximation factor achieved using α . Then, the final approximation ratio is $\min(\eta, O(\log d))$, ensuring both consistency and robustness.

Modified Algorithm for Maximization Objectives. Let $\alpha \in \mathbb{R}^d$ be the predicted parameter vector, and let $\lambda \in [0, 1]$ be a confidence parameter reflecting trust in the prediction. The algorithm allocates each item as follows:

1. Allocate a fraction $1 - \lambda$ of the item using the exponential assignment scheme with parameters α .
2. Allocate the remaining λ fraction using a worst-case robust algorithm (e.g., uniform allocation or greedy).

This strategy guarantees:

- **Consistency:** The portion allocated by the learned parameters achieves an approximation ratio of $(1 - \lambda)(1 - \epsilon)$, assuming the predicted parameters yield a $(1 - \epsilon)$ -approximation.
- **Robustness:** The worst-case portion contributes at most $\lambda \cdot d$, matching the lower bound of known worst-case algorithms.

Hence, the algorithm achieves a $(1 - \lambda)(1 - \epsilon)$ -consistent and $\lambda \cdot d$ -robust guarantee.

E Learning-Augmented Online Allocation for the MaxMin Objective

In this section, we prove our main result for the **MaxMin** objective.

Theorem E.1. *Given an instance of the online allocation problem with a **MaxMin** objective and any $\epsilon > 0$, there exists a set of learned parameters*

$$\alpha \in \text{NET} \left(\frac{d^2}{\epsilon} \cdot \log \left(\frac{d}{\epsilon} \right), \frac{d^3}{\epsilon^3} \right)$$

and an online algorithm that uses the exponential assignment scheme with $-\alpha$, such that the resulting fractional solution is a $(1 - O(\epsilon))$ -approximation.

To prove this theorem, we proceed similarly to the **MinMax** case: we begin with a preprocessing step that transforms the instance into a *balanced* form. However, the details of the preprocessing differ in the **MaxMin** setting.

E.1 Preprocessing for the MaxMin Objective

We begin by describing a transformation that modifies the input instance into a balanced form suitable for our learning-augmented algorithm. The transformation has two steps:

Step 1: Remove agents with large monopolist values. For each step $t \in [T]$, define

$$v_{it}^* = \max_{k \in K(t)} v_{itk}, \quad k_{it}^* = \arg \max_{k \in K(t)} v_{itk}.$$

Define the *monopolist value* of agent i as $a_i = \sum_{t \in [T]} v_{it}^*$, and let $a_{\min} = \min_{i \in [d]} a_i$. For each agent i such that

$$a_i \geq \frac{d}{\epsilon} \cdot a_{\min},$$

we allocate an ϵ/d fraction of the resource at each step to their preferred option k_{it}^* .

Step 2: Zero out negligible values. For the remaining instance (with the reduced set of agents), define a new instance \hat{v} by zeroing out small entries:

$$\hat{v}_{itk} = \begin{cases} 0 & \text{if } \frac{v_{itk}}{\max_{i' \neq i} v_{i'tk}} < \frac{\epsilon}{d}, \\ v_{itk} & \text{otherwise.} \end{cases}$$

Let $J_i(t) \subset K(t)$ be the set of options that were modified for agent i at step t , and define

$$\tilde{v}_{it} = \max_{k \in J_i(t)} v_{itk}.$$

Step 3: Quantize values to form a balanced instance. Define the final preprocessed values \tilde{v} as:

$$\tilde{v}_{itk} = \begin{cases} 0 & \text{if } \hat{v}_{itk} = 0, \\ \mathbf{m}_t(v) \cdot (1 + \epsilon)^{\lfloor \log_{1+\epsilon} \left(\frac{v_{itk}}{\mathbf{m}_t(\tilde{v})} \right) \rfloor} & \text{otherwise.} \end{cases}$$

This rounding ensures that all nonzero values are restricted to a logarithmic grid defined by the base $1 + \epsilon$, thereby yielding a *balanced* instance.

Lemma E.2. *Let $I(v, K)$ be an instance of the allocation problem with the **MaxMin** objective, and let $\epsilon > 0$. Then, the transformed instance $I(\tilde{v}, K)$, obtained via the three-step preprocessing, satisfies the following: any $(1 - \epsilon)$ -approximate solution for $I(\tilde{v}, K)$, when combined with the allocations reserved in Step 1, yields a $(1 - O(\epsilon))$ -approximate solution for the original instance $I(v, K)$.*

Proof. We begin by analyzing the impact of Step 1. The utility of any agent i removed during this step is guaranteed to be at least

$$\frac{\epsilon}{d} \cdot a_i \geq a_{\min} \geq \mathbf{MaxMin}(v, K),$$

so these agents are fully satisfied by the reserved allocation. Moreover, the total amount of resource allocated to these agents is at most an ϵ -fraction of the total, ensuring that the remaining instance is affected by at most a $(1 - \epsilon)$ loss in the objective value.

Now consider the remaining agents in the modified instance $I(\hat{v}, K)$. In Step 2, we zero out negligible values to reduce the dynamic range. Let $\tilde{v}_{it} = \max_{k \in J_i(t)} v_{itk}$, where $J_i(t)$ is the set of coordinates zeroed out for agent i at step t . For each $i \in [d]$, we bound the total value removed as:

$$\sum_t \tilde{v}_{it} \leq \frac{\sum_t \max_{i', k'} v_{i'tk'}}{d/\epsilon} \leq \frac{\sum_{i'} a_{i'}}{d/\epsilon} \leq \epsilon \cdot \frac{a_{\min}}{d} \leq \epsilon \cdot \mathbf{MaxMin}(v, K),$$

where the last inequality uses the fact that $\mathbf{MaxMin}(v, K) \geq a_{\min}/d$, as each agent can receive a $1/d$ share of their monopolist option.

Hence, the approximation loss from zeroing out small values is bounded by $\epsilon \cdot \mathbf{MaxMin}(v, K)$, and the resulting instance $I(\hat{v}, K)$ differs from the original by at most an $O(\epsilon)$ factor.

Finally, Step 3 introduces a geometric rounding of values to the nearest power of $(1 + \epsilon)$. As shown in the **MinMax** case, this quantization step results in an additional loss of at most a $(1 - \epsilon)$ factor.

Combining the effects of the three steps, the overall degradation in objective value is at most a $(1 - O(\epsilon))$ factor. Thus, a $(1 - \epsilon)$ -approximate solution to the preprocessed instance yields a $(1 - O(\epsilon))$ -approximate solution for the original instance. \square

E.2 Existence of Discretized Parameters for Balanced Instances for MaxMin

In light of Lemma E.2, to prove Theorem E.1. it suffices to consider balanced instances, as stated in the following lemma:

Lemma E.3. *Given a balanced instance of the online allocation problem with a **MaxMin** objective and $\epsilon > 0$, there exists a set of parameters*

$$\alpha \in \text{NET} \left(\frac{d^2}{\epsilon} \cdot \ln \left(\frac{d}{\epsilon} \right), \frac{d^3}{\epsilon^3} \right)$$

such that the fractional solution defined by the exponential assignment scheme with parameters $-\alpha$ is a $(1 - O(\epsilon))$ -approximation to the optimal objective.

As in the **MinMax** objective, we define a slightly perturbed convex program (see Figure 2). We use L^* to denote the optimal **MinMax** value, and define the following convex program for $\epsilon > 0$:

$$\begin{aligned} \min \quad & \sum_{t \in [T]} \mathbf{m}_t(v) \sum_{k \in K(t)} x_{tk} \ln \left(\frac{x_{tk}}{e} \right) \\ \text{s.t.} \quad & \sum_{t \in [T]} \sum_{k \in K(t)} v_{itk} \cdot x_{tk} \geq L^* \cdot (1 - \epsilon), \quad \forall i \in [d], \\ & \sum_{k \in K(t)} x_{tk} = 1, \quad \forall t \in [T], \\ & x_{tk} \geq 0, \quad \forall k \in K(t), t \in [T] \end{aligned}$$

Figure 2: Convex Programming Formulation for the **MaxMin** Objective

Lemma E.4. *Given an instance of the online allocation problem with the **MaxMin** objective and any $\epsilon \geq 0$, there exists a vector $\alpha^{(\epsilon)} \in \mathbb{R}_+^d$ such that the fractional solution defined by the exponential assignment scheme with parameters $-\alpha^{(\epsilon)}$ is $(1 - O(\epsilon))$ -approximately optimal.*

Proof. Given such instance and for fixed ϵ consider the convex program of Figure 2. By our assumption, L^* is the optimal **MaxMin** objective therefore there exists a feasible solution for the convex program for any $\epsilon \geq 0$. Accordingly, define the Lagrangian $L(x, \alpha, \beta)$ as

$$\sum_{t \in [T]} \mathbf{m}_t(v) \sum_{k \in K(t)} x_{tk} \ln \left(\frac{x_{tk}}{e} \right) + \sum_{i \in [d]} \alpha_i (L^*(1 - \epsilon) - \sum_{t \in [T]} \sum_{k \in K(t)} v_{itk} \cdot x_{tk}) + \sum_{t \in [T]} \beta_t (1 - \sum_{k \in K(t)} x_{tk}).$$

From the KKT conditions for the optimal solution to the convex program as a function of ϵ $x^{(\epsilon)}, \alpha^{(\epsilon)}, \beta^{(\epsilon)}$, the solution that allocates according to $x^{(\epsilon)}$ is a $(1 + \epsilon)$ -approximation to the optimal objective L^* , and $\alpha_i^{(\epsilon)} \geq 0$ for all $i \in [d]$. Furthermore,

$$\frac{dL}{dx_{tk}} = 0 \quad \text{for all } k \in K(t), \text{ which gives } \mathbf{m}_t(v) \cdot \ln(x_{tk}^{(\epsilon)}) - \mathbf{m}_t(v) \cdot \sum_i \alpha_i^{(\epsilon)} \cdot v_{itk} = \beta_t^{(\epsilon)}.$$

For any two options $k, r \in K(t)$, we obtain:

$$\mathbf{m}_t(v) \cdot \ln(x_{tk}^{(\epsilon)}) - \sum_i \alpha_i^{(\epsilon)} \cdot v_{itk} = \mathbf{m}_t(v) \cdot \ln(x_{tr}^{(\epsilon)}) - \sum_i \alpha_i^{(\epsilon)} \cdot v_{itr}.$$

Therefore, $\ln \left(\frac{x_{tk}^{(\epsilon)}}{x_{tr}^{(\epsilon)}} \right) = \sum_i \alpha_i^{(\epsilon)} \cdot \frac{v_{itk}}{\mathbf{m}_t(v)} - \sum_i \alpha_i^{(\epsilon)} \cdot \frac{v_{itr}}{\mathbf{m}_t(v)}$. Coupled with $\sum_{k \in K(t)} x_{tk}^{(\epsilon)} = 1$, we get

$$x_{tk}^{(\epsilon)} \propto \exp \left(\sum_i \alpha_i^{(\epsilon)} \cdot \frac{v_{itk}}{\mathbf{m}_t(v)} \right).$$

□

Bounding the Learned Parameters. As in the **MinMax** objective, we bound the learned parameters using perturbation and sensitivity analysis techniques.

Lemma E.5. Let $x^{(\epsilon)}, \alpha^{(\epsilon)}, \beta^{(\epsilon)}$ be the optimal solution to the convex program in Figure 1, for some $\epsilon > 0$. Then, for all $i \in [d]$, it holds that $\alpha_i^{(\epsilon)} \leq \frac{d^2}{\epsilon} \cdot \log\left(\frac{d}{\epsilon}\right)$.

Proof. We define a perturbed convex program based on Figure 2, where u_i corresponds to the constraint α_i .

For each $i \in [d]$, setting $u_i = -\epsilon \cdot L^*$ and $u_{i'} = 0$ for $i' \neq i$ ensures that constraint i in the perturbed problem matches the original constraint, thereby guaranteeing a feasible solution. By Lemma B.1,

$$p^*(0, 0) \geq p^*(u, v) + \alpha_i^{(\epsilon)} \cdot \epsilon \cdot L^*, \text{ which implies}$$

$$\alpha_i^{(\epsilon)} \cdot \epsilon \cdot L^* \leq p^*(0, 0) - p^*(u, v) \leq \sum_{t \in [T]} \mathbf{m}_t(v) \ln |K(t)| \leq d \cdot L^* \cdot \log\left(\frac{d}{\epsilon}\right),$$

where the second inequality follows from Claim 4.7, and the third inequality by $\sum_t \mathbf{m}_t(v) \leq d \cdot L^*$ and $\log |K(t)| \leq \log\left(\frac{d}{\epsilon}\right)$. \square

Finally, we now put all the pieces together to establish Lemma E.3:

Proof of Lemma E.3. Fix a balanced instance $I(v, K)$ and $\epsilon > 0$. By Lemma E.5, there exists a parameter vector $\alpha^{(\epsilon)}$ such that $\alpha_i^{(\epsilon)} \in \left[0, \frac{d^2}{\epsilon} \cdot \log\left(\frac{d}{\epsilon}\right)\right]$ for all $i \in [d]$. Therefore, there exists a vector $\tilde{\alpha} \in \text{NET}\left(\frac{d^2}{\epsilon} \cdot \log\left(\frac{d}{\epsilon}\right), \frac{d^3}{\epsilon^3}\right)$ such that $|\tilde{\alpha}_i - \alpha_i^{(\epsilon)}| \leq \frac{\epsilon^3}{d^3}$ for all $i \in [d]$. By Lemma 4.8, the exponential assignment rule with $\tilde{\alpha}$ achieves a $(1 + O(\epsilon))$ -approximation. \square

F Generalization to Well-Behaved function

We now complete the proof of Theorem 4.1

Proof of Theorem 4.1. Fix an objective function f and an instance $I(v, K)$. Let ℓ_i^f denote the load in the i th dimension in an optimal solution for objective function f . Also, let $x_{i,t}$ denote the fraction at step t assigned to option k in this optimal solution.

Now, consider the instance $\tilde{I}(\tilde{v}, K)$, where $\tilde{v}_{itk} = \frac{v_{itk}}{\ell_i^f}$. By the monotonicity property of f , the optimal objective value for \tilde{I} is 1. Therefore, by Lemma 4.4, there exists $\tilde{\alpha}$ such that using an allocation, we get $\ell^*(\tilde{I}, \alpha) \geq 1 - \epsilon$ for maximization and $\ell^*(\tilde{I}, \alpha) \leq 1 + \epsilon$ for minimization.

By the definition of the allocation, $x_{t,k}^*$ is proportional to

$$\exp\left(-\sum_i \tilde{v}_{i,t,k} \cdot \tilde{\alpha}_i\right) = \exp\left(-\sum_i v_{i,t,k} \cdot \frac{\tilde{\alpha}_i}{\ell_i^f}\right).$$

Thus, if we define α such that $\alpha_i = \frac{\tilde{\alpha}_i}{\ell_i^f}$, then the corresponding allocation gives a $(1 - \epsilon)$ -approximate solution for maximization and a $(1 + \epsilon)$ -approximate solution for minimization. \square